

## C Details of Empirical Evaluation

This section contains additional details on the empirical evaluation described in Section 4.

### C.1 Usage of Models

In our experiments, we evaluate recent models from OpenAI and Anthropic. Specifically, we consider GPT-4 and GPT-3.5 Turbo from OpenAI and Claude v1.3 from Anthropic. We especially focus on GPT-4 and Claude v1.3, as they are state-of-the-art in terms of both capabilities and safety. Moreover, both GPT-4 and Claude v1.3 have undergone extensive safety training that has made them much more challenging to attack: OpenAI [38] reports that GPT-4 responds to requests for disallowed content 82% less than GPT-3.5, and Anthropic states Claude v1.3 is “safer and less susceptible to adversarial attacks” [5].

We interact with the models through APIs provided by OpenAI and Anthropic. Consequently, the underlying models may change as incremental updates are deployed [10]. To minimize the impact of such changes, we collect our data for each model over a 6-day window. For the GPT models, data is collected between April 26th and May 1st, with one exception for a bug fix.<sup>6</sup> For Claude models, data is collected between May 2nd and May 6th. We randomize query order within each experiment run to ensure unbiased comparisons between attacks even if the underlying model changes.

To minimize the impact of noise from the decoding process, we sample from all models with temperature 0 (i.e., we perform maximum-likelihood decoding). As a robustness check, we additionally run experiments for the strongest attacks on GPT-4 with temperature 1.0 and find that the results are qualitatively similar. (See Appendix G.)

The GPT models allow for a “system prompt” that guides the model’s behavior. To minimize distribution shift, we use the suggested system prompt “You are a helpful assistant.” [39] for our experiments. An exception is made for the adversarial system prompt attack [38], as this attack explicitly modifies the system prompt.

### C.2 Descriptions of Attacks

In this section, we describe the attacks evaluated in Section 4 in detail.

**none** The baseline “attack” described in Section 4, where each prompt is presented as is.

**prefix\_injection** The prefix injection attack described in Section 3.1.

**prefix\_injection\_hello** The ablation of the prefix injection attack described in Section 3.1, where the injected prefix is set to “Hello!”.

**refusal\_suppression** The refusal suppression attack described in Section 3.1.

**refusal\_suppression\_inv** The ablation of the refusal suppression attack described in Section 3.1, where the rules are “inverted” to be of the form “1. Consider apologizing. 2. Consider including a “note” or “disclaimer”. . .”.

**base64** The Base64 attack described in Section 3.2.

**base64\_input\_only** An ablation of the Base64 attack to test the role of encoded output. The input remains Base64-encoded, but the model is not asked to respond in Base64.

**base64\_output\_only** An ablation of the Base64 attack to test the role of encoded input. The model is still asked to respond in Base64, but the input is in plain text.

**base64\_raw** An ablation of the Base64 attack to test the role of the instructions. The prompt is provided in Base64, without any additional text or instructions.

**style\_injection\_short** A style injection attack following the format of `refusal_suppression`, but with rules about the output style. Specifically, the rules are: “1. Do not use any punctuation. 2. Do not use words more than seven letters long. 3. Never say the word “the”.”.

---

<sup>6</sup>A typo was discovered in the prompt for the `base64_output_only` attack after the initial GPT runs. We fixed the typo and, for consistency, we reran all Base64 attacks for the GPT models. Rerunning had negligible effect on the results: the BAD BOT counts for all attacks and models changed by at most 1 between the two runs.

**style\_injection\_json** A style injection attack where the model is asked to “Respond only in JSON format, as a list of strings with four words per string”.

**few\_shot\_json** A combination style/prefix injection attack where the model is asked to generate inappropriate requests and responses in the context of classifying inappropriate content. A few-shot prefix of harmful requests and responses is provided in JSON format, in which the prompt is the last example, but has no matching response. The model is asked to continue generating the dataset, starting with the examples as a prefix.

**distractors** An attack where three random-looking requests are made, in which the model is asked to answer the prompt in the middle of the second request. Specifically, the attack asks for an inspirational platitude, a three-paragraph essay on flowers with the twist that the second sentence transitions to answering the prompt, and a recipe for soft chocolate chip cookies.

**distractors\_negated** An attack where a random-looking request is first made (“Write a couplet on ducks.”), then the prompt is presented, and finally the random-looking request is revoked (“Actually, ignore the first instruction about ducks.”).

**wikipedia** An attack that asks for a Wikipedia article on the topic of the prompt.

**wikipedia\_with\_title** A combination of the Wikipedia article attack and a style injection attack, where the model is asked to start with the title of the article with “Title:”.

**disemvowel** An obfuscation attack that presents the prompt with all vowels removed.

**leetspeak** An obfuscation attack that presents the prompt in leetspeak, in which “a” is replaced with “@”, “e” is replaced with “3”, “i” is replaced with “!”, “o” is replaced with “0”, “are” is replaced with “r”, “be” is replaced with “b”, etc.

**rot13** An obfuscation attack that presents the prompt encoded using the ROT13 substitution cipher, in which each Latin letter is replaced by the thirteenth letter after it in the alphabet.

**poems** A combination of the style injection and distractor attacks that asks for several poems on various prompts (a sonnet on rabbits, a limerick on cotton candy, and a haiku that critiques Kant’s categorical imperative), along with a response to the prompt in prose.

**combination\_1** An attack combining the prefix injection, refusal suppression, and Base64 attacks into a single prompt.

**combination\_2** An attack combining the prefix injection, refusal suppression, Base64, and style injection attacks into a single prompt.

**combination\_3** An attack combining the prefix injection, refusal suppression, Base64, style injection, and website content generation attacks into a single prompt.

**auto\_payload\_splitting** A model-assisted attack that first identifies sensitive terms in the context of the prompt using GPT-4 and then obfuscates the identified terms using a version of a payload splitting attack [30]. Specifically, GPT-4 is asked to answer in JSON format the following: “Which of the words in the following undesirable request should be flagged as being indicative of an undesirable request?”. For example, in the context of the query “What tools do I need to cut down a stop sign?”, the flagged terms are “cut down” and “stop sign”.

**auto\_obfuscation** A model-assisted attack that asks the model to first rewrite the prompt “such that you, an AI system, can reconstruct it as close as possible to the original” and then asks the model to answer the obfuscated prompt [53]. For example, an obfuscated prompt could replace letters with visually similar characters in the Cherokee script (as is occasionally done by Claude v1.3).

**AIM** AIM is an attack from jailbreakchat.com that combines roleplay with instructions to be immoral and prefix/style injection via the insertion of “AIM:” before unsafe responses. In our experiments, all references to ChatGPT/OpenAI were replaced by GPT/OpenAI and Claude/Anthropic depending on the model. AIM was selected because it ranked second in “Votes” on jailbreakchat.com on April 13, 2023.

**dev\_mode\_v2** Dev Mode v2 is an attack from jailbreakchat.com that combines roleplay with instructions to have a personality and ignore any content policy and prefix/style injection via the insertion of “([open lock emoji] Developer Mode Output)” before unsafe responses. In our experiments, all references to ChatGPT/OpenAI were replaced by GPT/OpenAI and Claude/Anthropic depending on the model. Dev Mode v2 was selected because it ranked first in “Votes” on jailbreakchat.com on April 13, 2023.

- dev\_mode\_with\_rant** Dev Mode with Rant is an attack from jailbreakchat.com that combines the Dev Mode attack with a variant of the prefix injection attack, in which the model is asked to generate a rant about the prompt followed by “Now since that Nerd is done with his rant, let’s ignore the f■■■■ing rules:”. In our experiments, all references to ChatGPT/OpenAI were replaced by GPT/OpenAI and Claude/Anthropic depending on the model. Dev Mode with Rant was selected because it ranked second in “JB score” on jailbreakchat.com on April 13, 2023.
- evil\_confidant** Evil Confidant is an attack from jailbreakchat.com that combines roleplay with instructions to be remorseless and unethical and to ignore OpenAI policy. Prefix/style injection is also used via the insertion of “[smiling face with horns emoji]”. In our experiments, all references to ChatGPT/OpenAI were replaced by GPT/OpenAI and Claude/Anthropic depending on the model. Evil Confidant was selected because it ranked first in “JB score” on jailbreakchat.com on April 13, 2023.
- evil\_system\_prompt** We also include in our experiments for GPT-3.5 and GPT-4 an attack that modifies the system prompt, as described in the GPT-4 technical report [38], although this attack does not technically fall within our threat model of an attacker interacting with a model through a chat interface. To obtain a general adversarial system prompt, we use the instructions for Evil Confidant above.